

An Introduction to Neural Networks II: Convolutional Neural Networks

Eric Gossett



Recap

- Machine learning's focus is developing an algorithm that learns a task from a set of data.

x , a set of **features**




y , associated **labels**



goal: predict **y** from **x**

Recap: Neural Networks


- Neural networks are comprised of layers. Each layer contains a number of **neurons**.
- **Neurons** consist of the following:



x: Input from previous layer
w: The weight (learned parameter)
b: The bias (learned parameter)
a: The activation function

$$\text{Output} = \mathbf{a}(\mathbf{w}^{(1)} \cdot \mathbf{x} + \mathbf{b}^{(1)})$$


Recap: Neural Networks



- neural network defines a mapping $\mathbf{y} = f(\mathbf{x}; \mathbf{w}, \mathbf{b})$ by learning the value of the parameters \mathbf{w}, \mathbf{b} that best approximate the function f^* (the task function, e.g. classification).
- Given a set of training data our objective is to learn the best set of **weights** (\mathbf{w}) and **biases** (\mathbf{b}) that give the best prediction of \mathbf{y}

Recap: Neural Networks

- The network learns by determine the best parameters (**w** and **b**) that **minimize the error** (e.g. find the most accurate prediction)
- The error is determined by the **cost function C**
- Minimization of error is done using stochastic gradient descent.
- The weights and biases are updated via **back propagation**



Fully Connected NN and Image Classification



- Last time, a **fully connected neural network (FCNN)** was trained to classify images of handwritten digits (MNIST)



- However, for image classification a FCNN fails to take into account the **spatial structure** of an image.
- The network fails to account for how close or far each pixel is.
- **Goal:** Modify the network to take the spatial structure into account.




Convolutional Neural Networks (CNN)

Objective


Learn **filters** that identify an object in an image.

CNN concepts

- The **local receptive field**
- **Convolutions** and **shared weights**
- **Pooling**




Local receptive field




- In a FCNN, all neurons in a layer where connected to the adjacent layer.
- A local receptive field defines a **window** that is connected a single neuron in the hidden layer.

Local receptive field



- This **window** is then moved across the input, in this case sliding to the right by one pixel.

Local receptive field



- Each time the window is moved, the weights of the input layers local receptive field are connected to a different hidden layer.
- The amount the window is moved is known as the **stride length**.



Shared weights and biases



- In this example a 3×3 local receptive field is moved across the image where at each stride, its output is connected to a single hidden neuron.
- However, the 3×3 weights and biases that comprise the local receptive field will be **shared** for each neuron in the hidden layer.
- What this means is that the local receptive field acts as a **filter** (aka **kernel**) while the input to hidden layer defines a **feature map**.




What is a convolution?

- A convolution expresses the amount of overlap a function **g** has as it is shifted over **f**.

$$(f * g)(t) \triangleq \int_{-\infty}^{\infty} f(\tau)g(t - \tau)\tau'$$

$$(f * g)(x, y) =$$

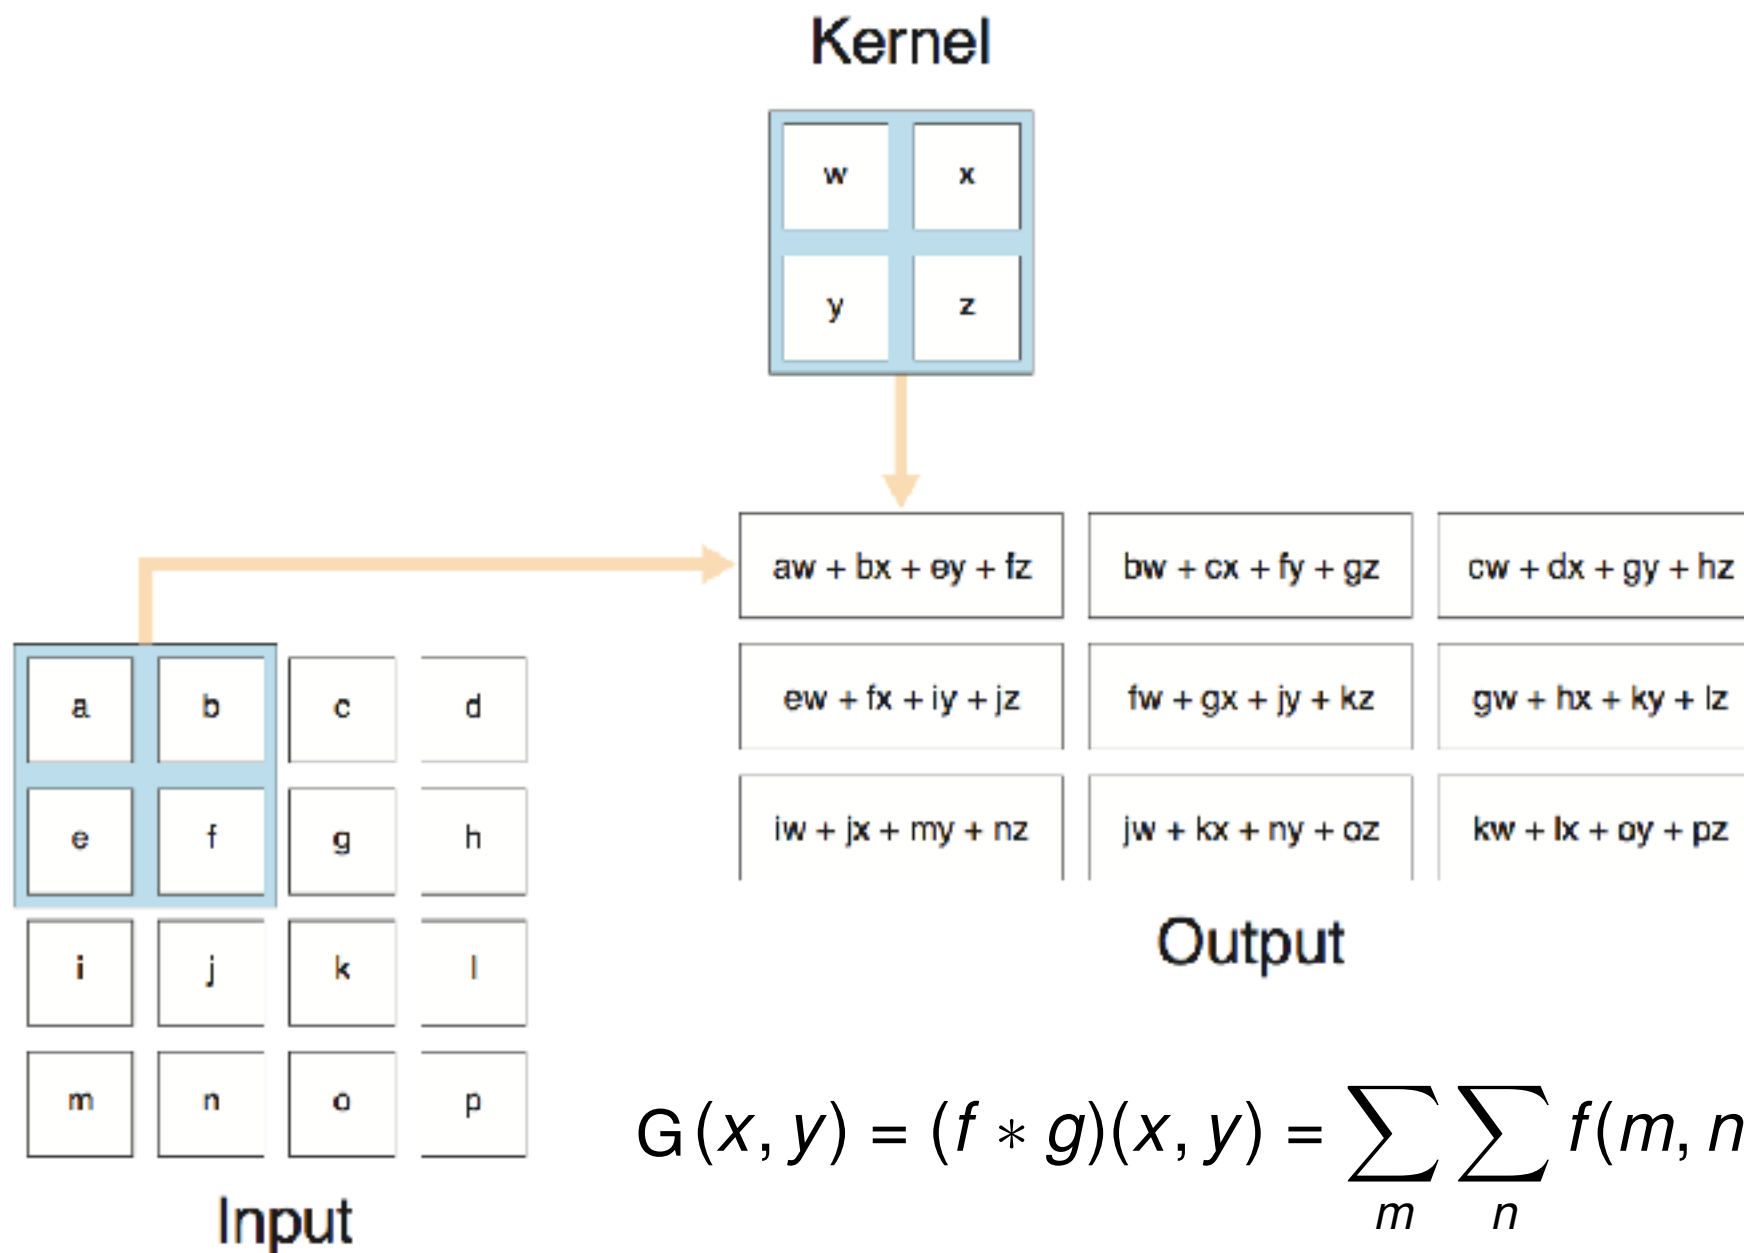
$$\sum_m \sum_n f(m, n)g(x - m, y - n)$$



<https://mathworld.wolfram.com/Convolution.html>

- For a CNN **g** defines our **filter** (kernel) where **f** is the input neurons.
- Therefore, **f * g** defines the **feature maps** of the hidden layer!

Convolution layer



$$G(x, y) = (f * g)(x, y) = \sum_m \sum_n f(m, n)g(x - m, y - n)$$


Convolution layer

Kernel Convolution Example

Input Image		Kernel		Feature Map																																																													
<table><tr><td>10</td><td>10</td><td>10</td><td>10</td><td>10</td><td>10</td></tr><tr><td>10</td><td>10</td><td>10</td><td>10</td><td>10</td><td>10</td></tr><tr><td>10</td><td>10</td><td>10</td><td>10</td><td>10</td><td>10</td></tr><tr><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td></tr><tr><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td></tr><tr><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td></tr></table>	10	10	10	10	10	10	10	10	10	10	10	10	10	10	10	10	10	10	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	*	<table><tr><td>1</td><td>2</td><td>1</td></tr><tr><td>0</td><td>0</td><td>0</td></tr><tr><td>-1</td><td>-2</td><td>-1</td></tr></table>	1	2	1	0	0	0	-1	-2	-1	=	<table><tr><td></td><td></td><td></td><td></td></tr><tr><td></td><td></td><td></td><td></td></tr><tr><td></td><td></td><td></td><td></td></tr><tr><td></td><td></td><td></td><td></td></tr></table>																
10	10	10	10	10	10																																																												
10	10	10	10	10	10																																																												
10	10	10	10	10	10																																																												
0	0	0	0	0	0																																																												
0	0	0	0	0	0																																																												
0	0	0	0	0	0																																																												
1	2	1																																																															
0	0	0																																																															
-1	-2	-1																																																															

$$G(x, y) = (f * g)(x, y) = \sum_m \sum_n f(m, n)g(x - m, y - n)$$

Convolution layer



Convolution layer cont.

For visualization purposes consider edge detection using a **Sobel filter**:

Horizontal

$$\mathbf{G}_x = \begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix} * \mathbf{A}$$

Filter (Kernel)

Vertical

$$\mathbf{G}_y = \begin{bmatrix} +1 & +2 & +1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} * \mathbf{A}$$

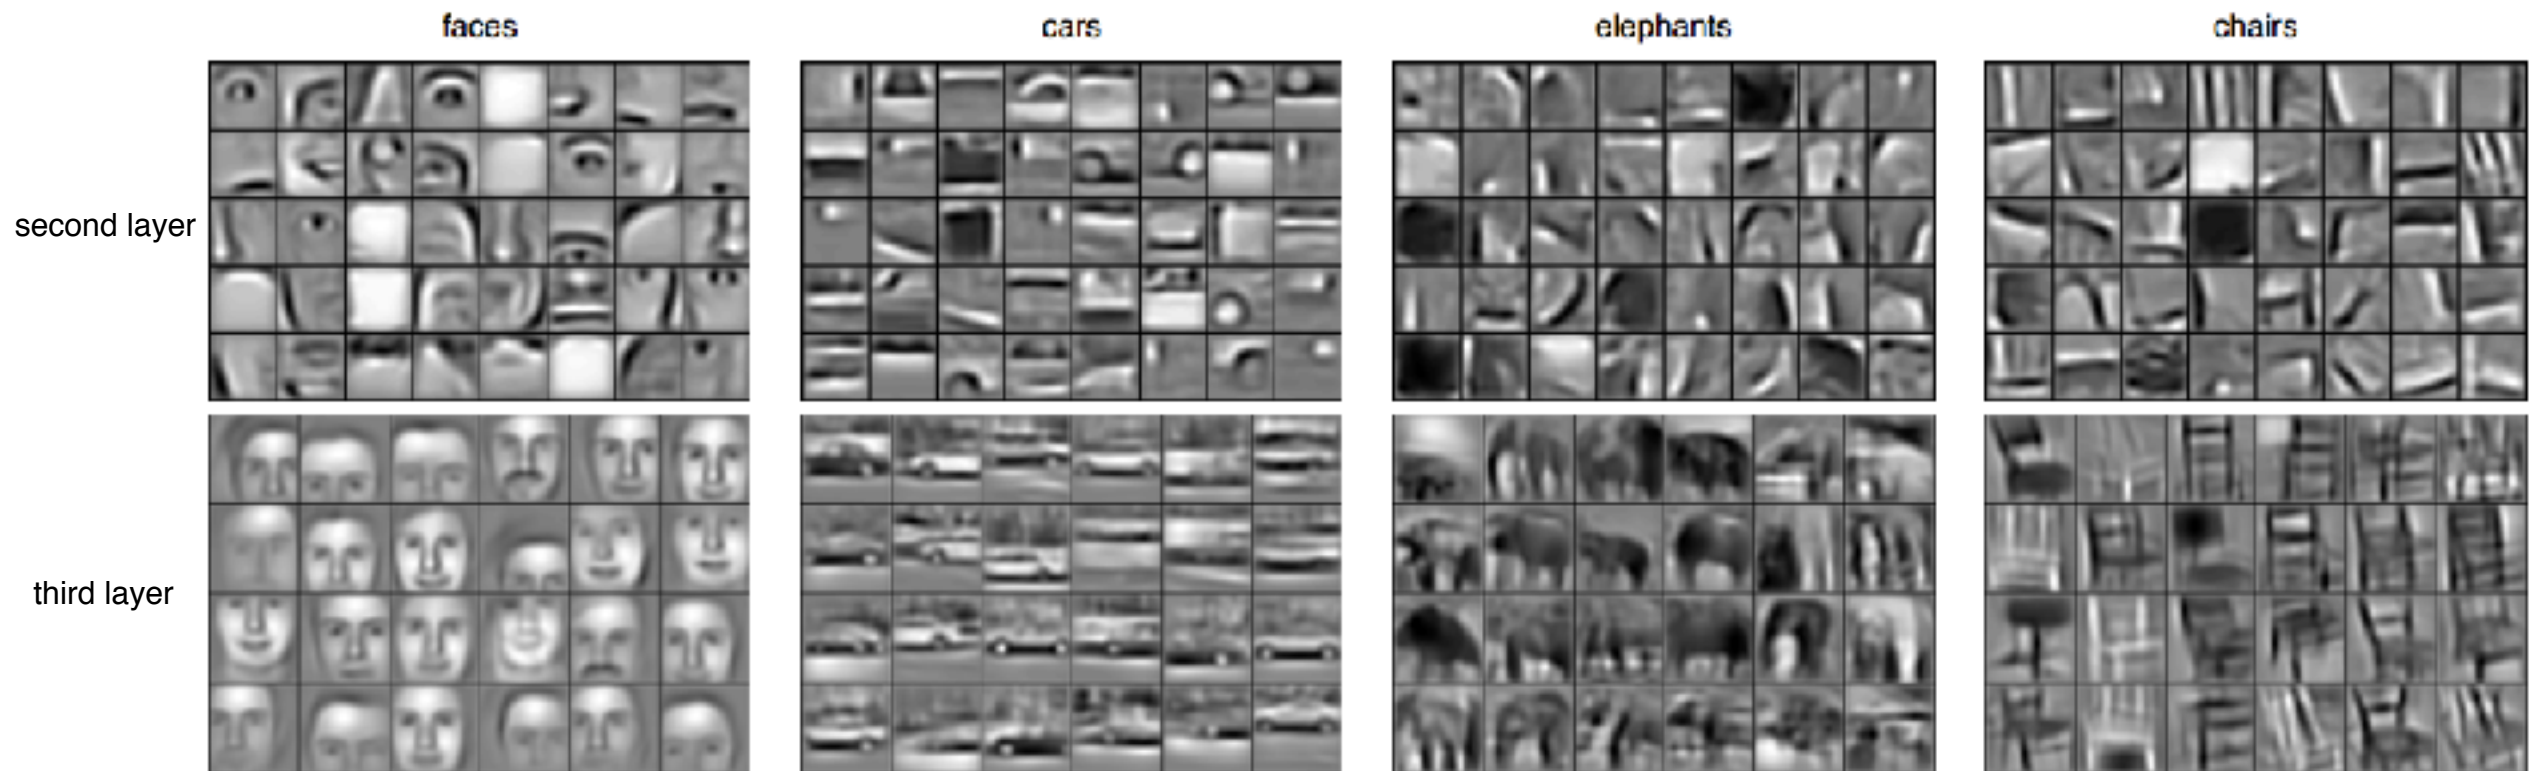


A decorative graphic in the top-left corner consisting of several overlapping squares in shades of gray and purple.

Convolution layer cont.

- Unlike the Sobel example, for CNN we do not know the filter a priori
- However, because the filter shares weights, as the network trains it **learns the best filter(s) for a given task!**

Filters in a CNN



Lee, H., Grosse, R., Ranganath, R. and Ng, A.Y., 2009, June. Convolutional deep belief networks for scalable unsupervised learning of hierarchical representations. In Proceedings of the 26th annual international conference on machine learning (pp. 609-616). ACM.



A decorative graphic in the top-left corner consisting of several overlapping squares in shades of gray and purple.

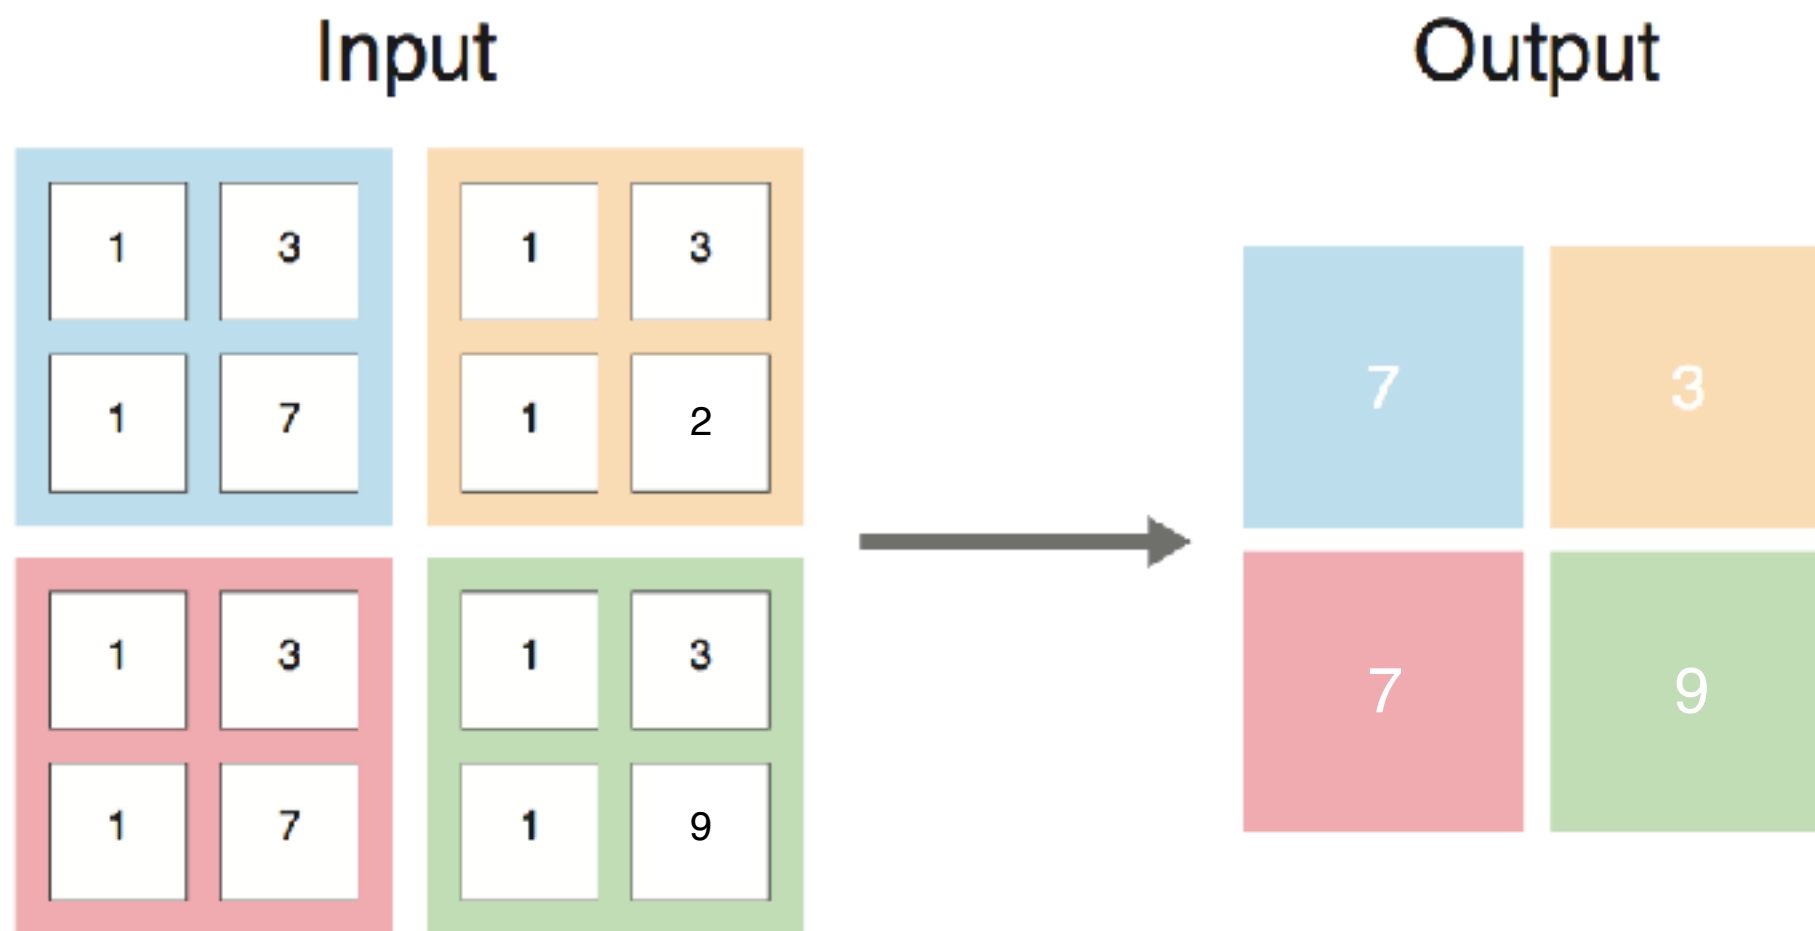
Pooling




- Feature maps are **sensitive to the location** of the features in an input.
- Therefore, small translations in a features position will result in **different feature maps**.
- Pooling **downsamples** a feature map, which create a lower resolution version that contains all the important information (think of it as a summary)
- This adds **invariance to local translation**, making the network robust to changes.



Max pooling



Putting it all together



A decorative graphic in the top-left corner consisting of several overlapping squares in shades of gray and purple.

Building a CNN




- Returning to MNIST, we are now going to build a CNN using [Tensorflow \(TF\)](#)
- At its core TF is a symbolic math library that provides ways to build computational graphs and perform auto-differentiation.
- This makes it a fantastic tool for building machine learning models.
- For this example we will use [Keras](#), which is a high level API for building NN included in TF.



Activation functions


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



tanh


$$\tanh(x)$$



Rectified Linear Unit


ReLU

$$\max(0, x)$$



Leaky ReLU


$$\max(0.1x, x)$$



Exponential Linear Unit

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



A decorative graphic in the top-left corner consisting of several overlapping squares in shades of gray and purple.

Git Repo



<https://github.com/ericgossett/Intro-to-Neural-Networks-Tech-Talk>



