

An Introduction to Neural Networks and Deep Learning

Eric Gossett



A decorative graphic in the top-left corner consisting of several overlapping squares in shades of gray and purple.

Objectives




- Provide an introduction to deep learning with an emphasis on first principles:
 - Intro to machine learning
 - The theory of neural networks
 - How to create a neural network from scratch.
- You will learn algorithms that use data to **automate tasks** and **make predictions**
- Can apply these techniques to numerous tasks such as: image classification, voice recognition, object detection, text translation etc.



What is Deep Learning?

- Deep learning is **not new!**
(neural networks [NN] were introduced in the 1950's)
- Resurgence is due to improvements of learning algorithms, more computational power and amazing performance metrics.
- In particular the spark was the **ImageNet** Challenge where NN outperformed all previous methods!



Machine learning (ML)

- ML focus is developing an algorithm that learns a task from a set of data.

x, a set of **features**



y, associated **labels**


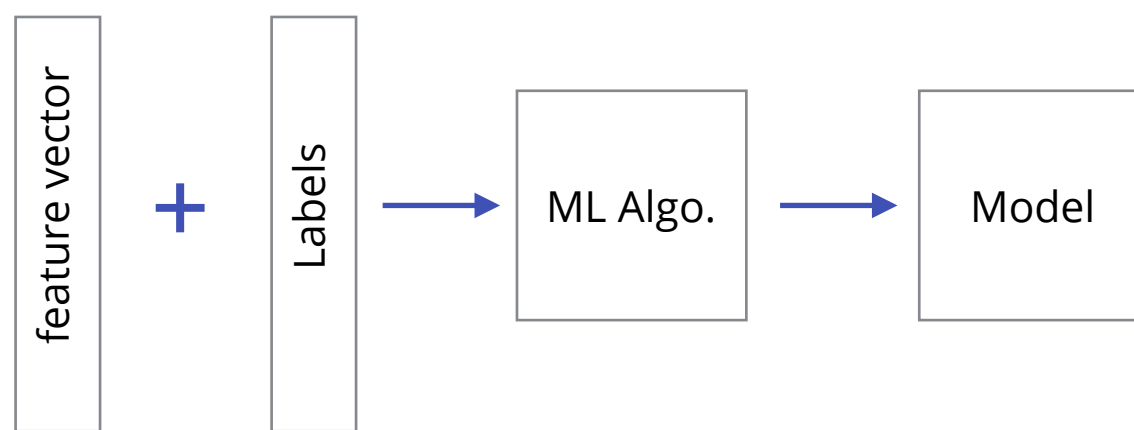


goal: predict **y** from **x**

Machine learning (ML)


- Two types of tasks **regression** and **classification**.
- **unsupervised learning** - Infer a model from unlabeled data.
- **supervised learning** - Infer a models that maps **feature** to a **label** by training on a labeled set of data.

Training set




Unsupervised Learning Methods

Clustering



Hidden Markov Model




Single Value Decomposition (SVD)


$$A = U\Sigma V^T$$

Supervised Learning methods

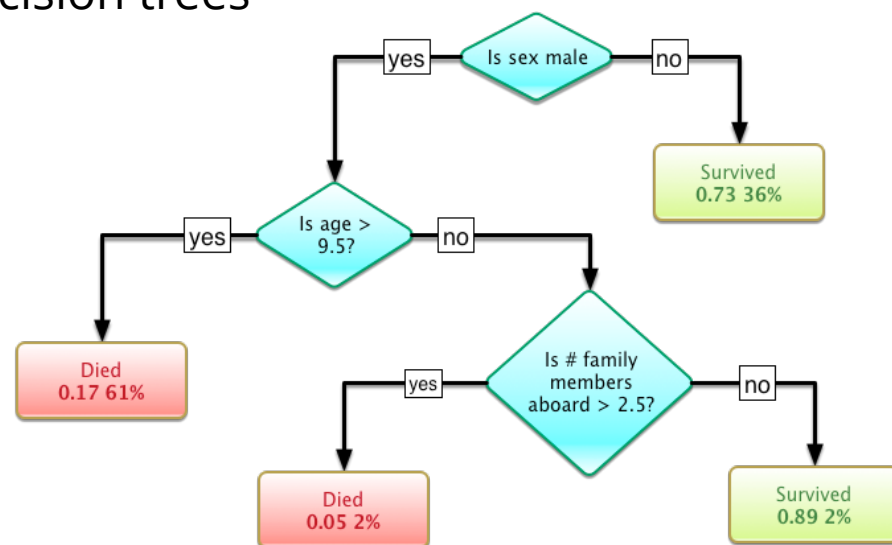
Linear regression




Support vector machines (SVM)



Decision trees




Neural networks (NN)




Training set

- ML algorithms require a training set to learn parameters


A set of **feature vectors**





with their associated **labels**



Given a **new** feature vector predict an **unknown** y



From training data to prediction (learning parameters)




Learn the parameters of the network:

w - the weights
b - bias

Classic Neural Network


- Neural networks are comprised of layers. Each layer contains a number of **neurons**.
- In a traditional NN **neurons** have the following form:



Neural Networks



- Neural networks aim to approximate some a task function f^* by composing together many different functions $f(\mathbf{x}) = f^{(3)}(f^{(2)}(f^{(1)}(\mathbf{x})))$.



- neural network defines a mapping $\mathbf{y} = f(\mathbf{x}; \mathbf{w})$ by learning the value of the parameters \mathbf{w} that best approximate the function f^* .






Learning the weights and biases



- Given a set of training data our objective is to learn the best set of **weights (\mathbf{w})** and **biases (\mathbf{b})** that give the best prediction of **\mathbf{y}**
- This is an **variational problem**: Determine the best parameters (**\mathbf{w}** and **\mathbf{b}**) that **minimize the error** (e.g. find the most accurate prediction)
- Learning is done in the following steps:
 - Feed forward
 - Back propagation of error
 - Gradient descent




Neural Networks



$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \quad \mathbf{w} = \begin{bmatrix} w_{11} & w_{21} & \cdots & w_{n1} \\ w_{12} & w_{22} & \cdots & w_{n2} \\ \vdots & \vdots & \ddots & \vdots \\ w_{1m} & w_{2m} & \cdots & w_{nm} \end{bmatrix} \quad \mathbf{b} = \begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_n \end{bmatrix}$$

NN: feed forward



$$\mathbf{z}^{(1)} = \mathbf{w}^{(1)} \cdot \mathbf{x} + \mathbf{b}^{(1)} = \begin{bmatrix} w_{11} & w_{21} & w_{31} \\ w_{12} & w_{22} & w_{32} \\ w_{13} & w_{23} & w_{33} \\ w_{14} & w_{24} & w_{34} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} + \begin{bmatrix} b_1 \\ b_2 \\ b_3 \\ b_4 \end{bmatrix}$$

$$= \begin{bmatrix} (w_{11} \cdot x_1 + w_{21} \cdot x_2 + w_{31} \cdot x_3) + b_1 \\ (w_{12} \cdot x_1 + w_{22} \cdot x_2 + w_{32} \cdot x_3) + b_2 \\ (w_{13} \cdot x_1 + w_{23} \cdot x_2 + w_{33} \cdot x_3) + b_3 \\ (w_{14} \cdot x_1 + w_{24} \cdot x_2 + w_{34} \cdot x_3) + b_4 \end{bmatrix}$$



Feed forward: Activation function




- During learning want small changes in **w** or **b** to result in small changes to **z** (the output).
- For a traditional neuron this is not the case, since a small change in either can flip the neuron.
- Therefore, must pass **z** to a special function known as the activation function. It has the following properties:
 - Has a derivative that can be computed
 - Is non-decreasing
 - Has horizontal asymptotes at 0 and 1 (or -1 and 1)



Activation functions


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



tanh


$$\tanh(x)$$



Rectified Linear Unit


ReLU

$$\max(0, x)$$



Leaky ReLU


$$\max(0.1x, x)$$




Exponential Linear Unit

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$





NN: feed forward



activation function **Sigmoid**


$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

$$\mathbf{z}^{(1)} = \mathbf{w}^{(1)} \cdot \mathbf{x} + \mathbf{b}^{(1)}$$



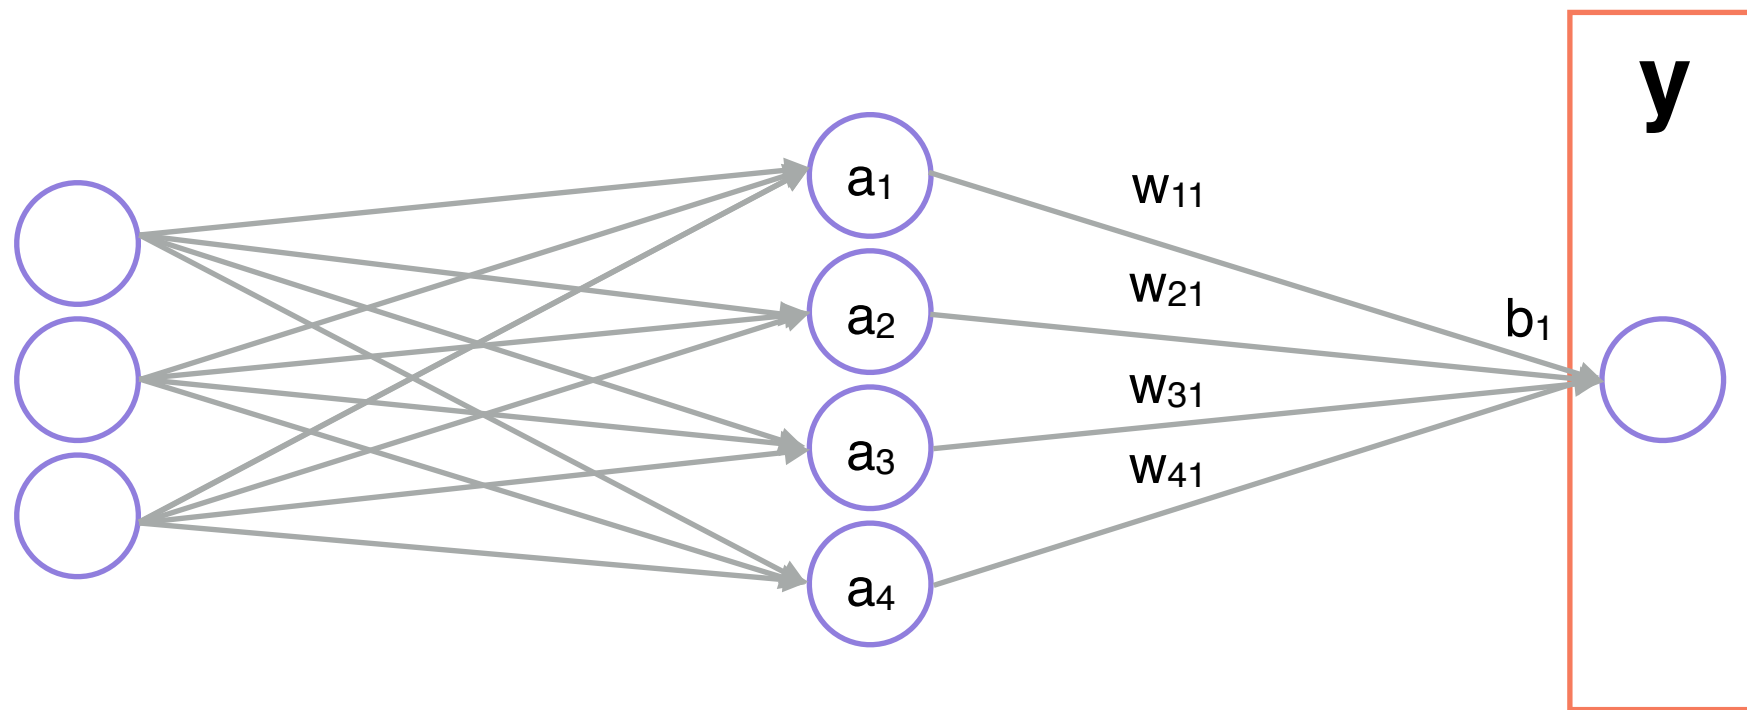
$$\mathbf{w} + \Delta \mathbf{w} \rightarrow \mathbf{y} + \Delta \mathbf{y}$$

NN: feed forward



$$\mathbf{a} = \sigma(\mathbf{w}^{(1)} \cdot \mathbf{x} + \mathbf{b}^{(1)}) = \sigma(\mathbf{z}^{(1)}) = \begin{bmatrix} a_1 \\ a_2 \\ a_3 \\ a_4 \end{bmatrix}$$

NN: feed forward



$$\mathbf{y} = \sigma(\mathbf{w}^{(2)} \cdot \mathbf{a} + \mathbf{b}^{(2)})$$

$$= \sigma\left(\begin{bmatrix} w_{11} & w_{21} & w_{31} & w_{41} \end{bmatrix} \begin{bmatrix} a_1 \\ a_2 \\ a_3 \\ a_4 \end{bmatrix} + b_1\right)$$



■ NN: Determining prediction accuracy



- Now we have a predicted value, how do we determine how good it is?
- This is known as the cost function. Many forms exist however the simplest is the mean squared error (MSE)

$$C(w, b) = \frac{1}{2n} \sum_i ||y_i - \tilde{y}_i||_2^2 = \frac{1}{2n} \sum_x ||\mathbf{y}(x) - \mathbf{a}^L(x)||^2 = \frac{1}{n} \sum_x C_x$$



Minimization of error




- Want to update the weights and biases in order to drive the network towards the $f(x)$ that best approximates f^*
- The metric to calculate the error in our network is the **cost function**

$$C(w, b) = \frac{1}{2n} \sum_i ||y_i - \tilde{y}_i||_2^2 = \frac{1}{2n} \sum_x ||\mathbf{y}(x) - \mathbf{a}^L(x)||^2 = \frac{1}{n} \sum_x C_x$$

- Want to minimize the cost (error) by changing the weights and biases.
- Therefore, we need two things: 1) way to update our weights and biases based on the error. 2) way to minimize the cost (error).



Gradient Descent



<https://saugatbhattarai.com.np/what-is-gradient-descent-in-machine-learning/>

- Recall this is a **variational problem** in which we want to pick the best \mathbf{w} and \mathbf{b} such that we minimize the error.
- In other words we want to find the minimum which is done using **gradient descent**



Gradient Descent caveats

1. Before we calculate the gradient we need some way to relate the error to the \mathbf{w} and \mathbf{b} such that we can update them based off the gradient
2. Calculating the total gradient for the entire feature space is expensive! This will require every entry of the training set (yikes)

Solutions

1. Use something known as **back propagation** to update the \mathbf{w} and \mathbf{b}
2. Use **stochastic gradient descent** to approximate the gradient from a smaller random batch



NN: Back propagation

- Want to minimize the cost (error) by changing the weights and biases.

$$C = \frac{1}{2n} \sum_i \|y_i - \tilde{y}_i\|_2^2 = \frac{1}{2n} \sum_x \|\mathbf{y}(x) - \mathbf{a}^L(x)\|^2 = \frac{1}{n} \sum_x C_x$$

- Want to back propagate error to update the weight and bias.
- Minimize \mathbf{C} using gradient descent.
- How do we calculate the gradient? Also how do we update the weights?
- Chain rule to the rescue!
- **Starting point:** the error for some input in the network is defined by:

$$\delta_j^l = \frac{\partial C}{\partial z_j^l}.$$

Back propagation

- Derive an expression for the error at the final layer:

$$\delta_j^L = \sum_k \frac{\partial C}{\partial a_k^L} \frac{\partial a_k^L}{\partial z_j^L} \quad \xrightarrow{\text{When } j=k \text{ else } 0} \quad \delta_j^L = \frac{\partial C}{\partial a_j^L} \frac{\partial a_j^L}{\partial z_j^L}$$
$$\delta_j^L = \frac{\partial C}{\partial a_j^L} \sigma'(z_j^L).$$
$$\frac{\partial C}{\partial a_j^L} = a_j^L - y_j$$
$$\sigma'(z) = \sigma(z)(1 - \sigma(z)).$$

In Matrix notation:

$$\delta^L = (\mathbf{a}^L - \mathbf{y}) \odot \sigma(\mathbf{z}^L)$$

Back propagation

- Derive an expression for the error at all previous layers:

$$\delta_j^l = \sum_k \frac{\partial C}{\partial z_k^{l+1}} \frac{\partial z_k^{l+1}}{\partial z_j^l} \quad \rightarrow \quad \delta_j^l = \sum_k \delta_j^{l+1} \frac{\partial z_k^{l+1}}{\partial z_j^l}.$$

$$\begin{aligned} \frac{\partial z_k^{l+1}}{\partial z_j^l} &= \frac{\partial}{\partial z_j^l} \left(\sum_j w_{kj}^{l+1} \sigma(z_j^l) + b_k^{l+1} \right) \\ &= w_{kj}^{l+1} \sigma'(z_j^l). \end{aligned}$$

In Matrix notation:

$$\delta^l = [(\mathbf{w}^{l+1})^T \delta^{l+1}] \odot \sigma(\mathbf{z}^l)$$

Back propagation computing gradients

- Finally we can compute the gradients!

$$\frac{\partial C}{\partial b_j^l} = \frac{\partial C}{\partial z_j^l} \frac{\partial z_j^l}{\partial b_j^l} \quad \leftarrow \begin{array}{l} \text{Term equals to 1} \\ \text{Also recall:} \\ \delta_j^l = \frac{\partial C}{\partial z_j^l} \end{array}$$

$$\frac{\partial C}{\partial b_j^l} = \delta_j^l$$

$$\frac{\partial C}{\partial w_{jk}^l} = \frac{\partial C}{\partial z_k^l} \frac{\partial z_k^l}{\partial w_{jk}^l}$$

$$\frac{\partial C}{\partial w_{jk}^l} = \delta_j^l a_k^{l-1}$$

In Matrix notation:

$$\frac{\partial C}{\partial \mathbf{b}} = \boldsymbol{\delta}$$

$$\frac{\partial C}{\partial \mathbf{w}} = \boldsymbol{\delta}^l \mathbf{a}^{l-1}$$

NN: Learning via Gradient Descent.


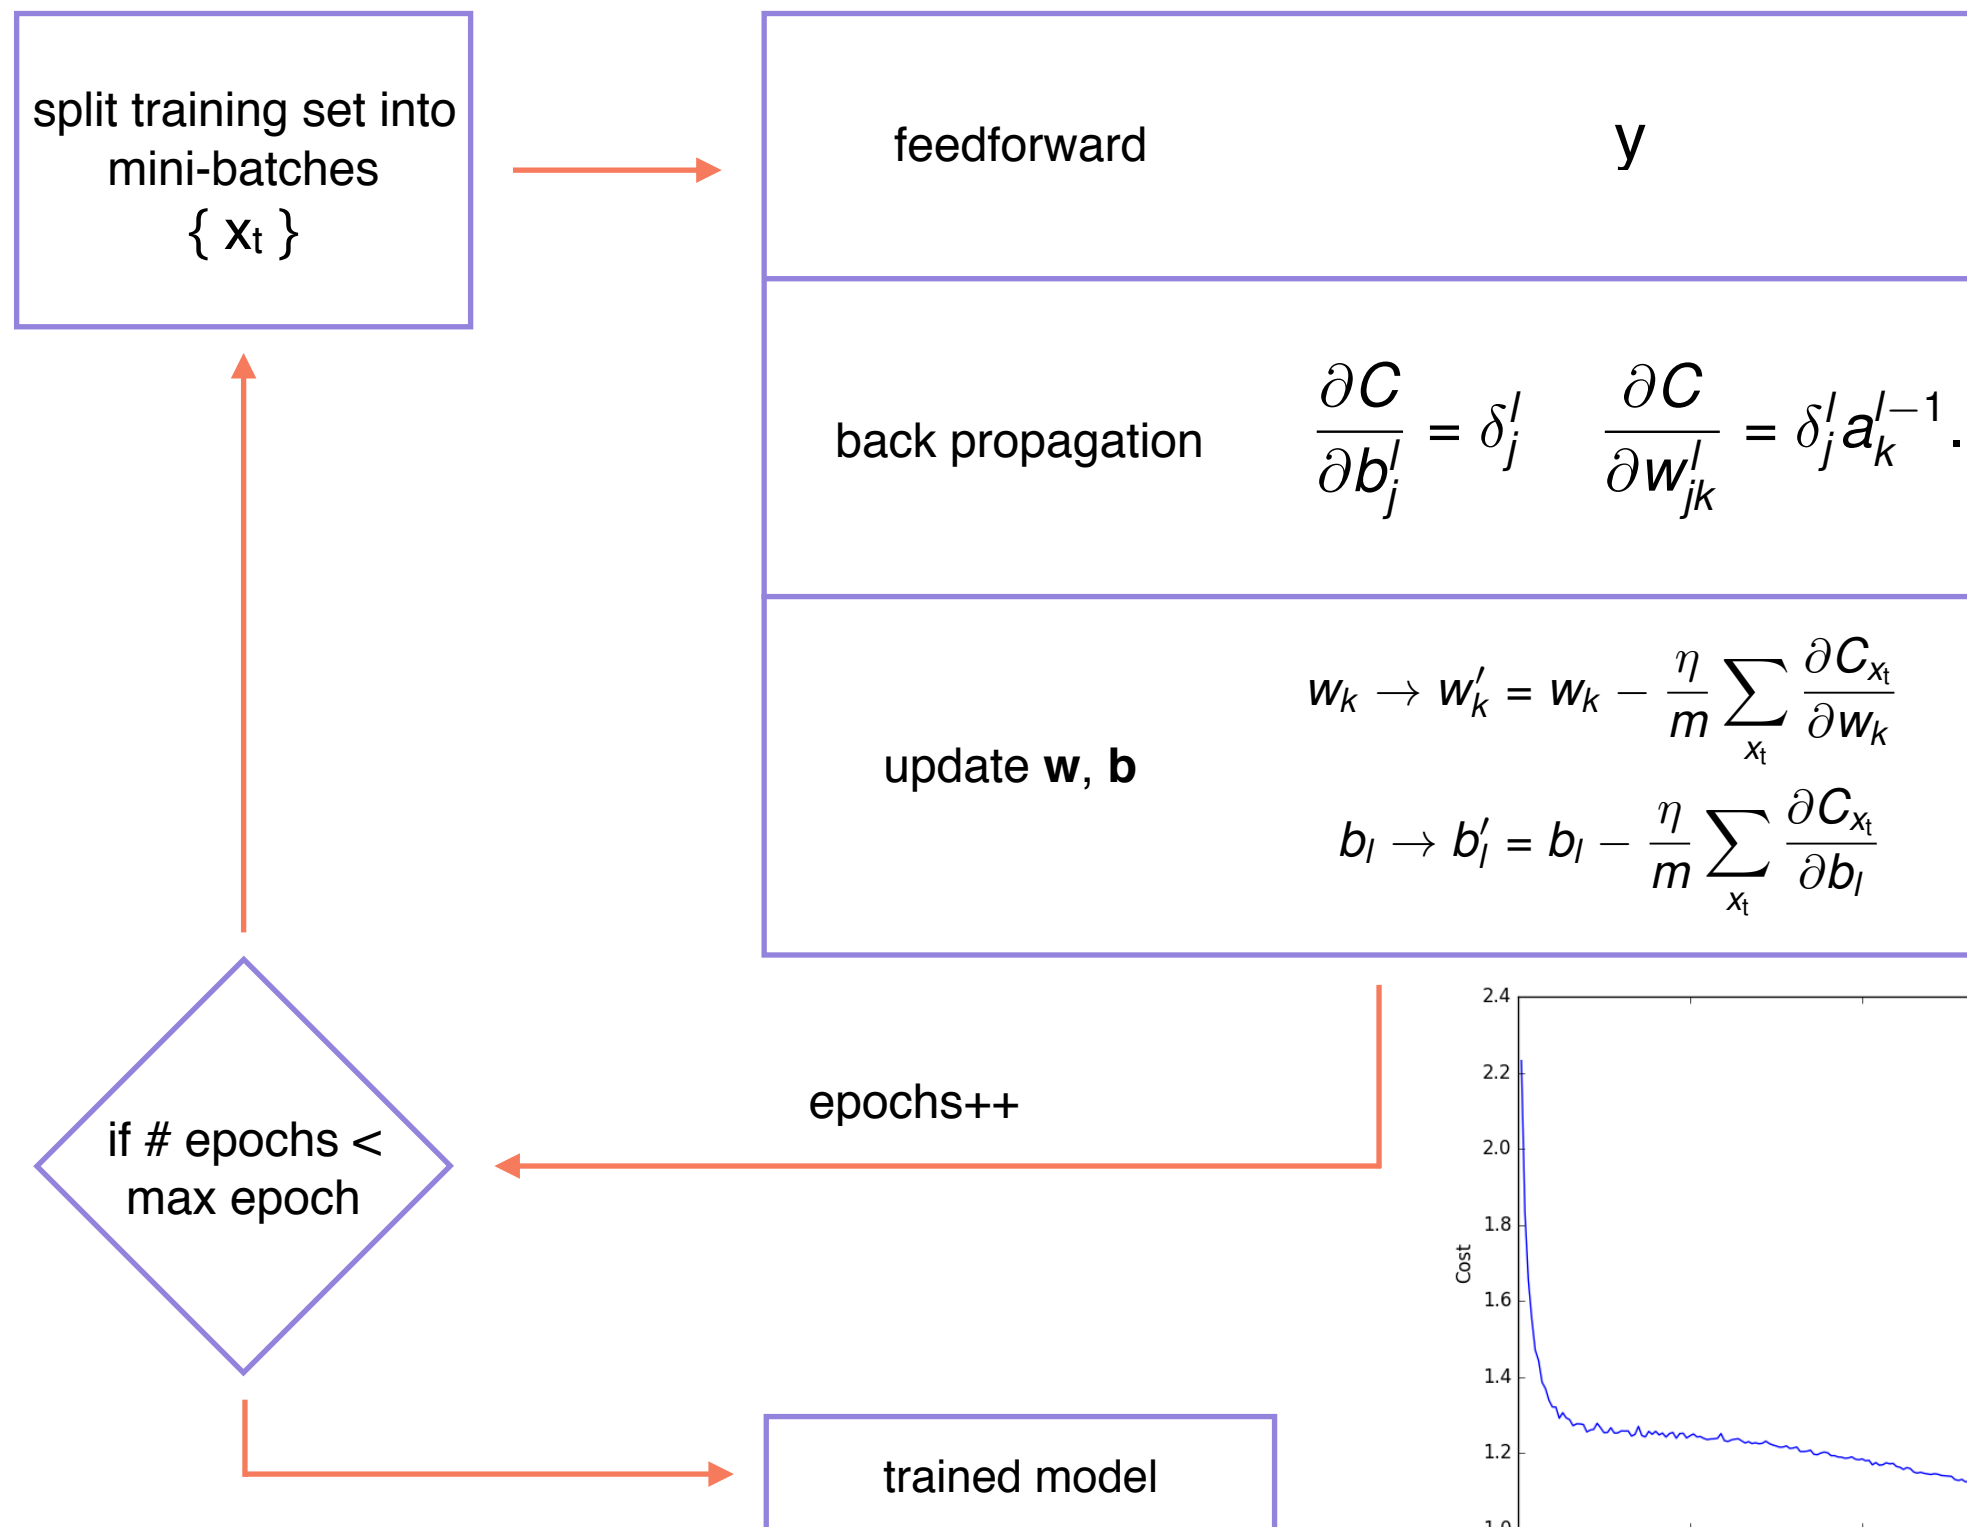
- Calculating the total gradient is expensive, so must approximate.
- Use stochastic gradient descent to break training set into **mini-batches** $\{x_t\}$ of size m , such that the gradient is:

$$\nabla C = \frac{1}{n} \sum_x \nabla C_x \approx \frac{1}{m} \sum_{x_t} \nabla C_{x_t}$$

- For each mini-batch, the components of the gradient are calculated using **back propagation**.
- Weights/biases updated via:


$$w_k \rightarrow w'_k = w_k - \frac{\eta}{m} \sum_{x_t} \frac{\partial C_{x_t}}{\partial w_k}$$
$$b_l \rightarrow b'_l = b_l - \frac{\eta}{m} \sum_{x_t} \frac{\partial C_{x_t}}{\partial b_l}$$

NN: Learning via Gradient Descent.



NN: Learning via Gradient Descent.

- The entire process is known as an **epoch**.
- This is then repeated over **multiple epochs** until the cost reaches a minimum.



Building a NN: Training set


- **MNIST:** Modified National Institute of Standards and Technology is a collection of hand written digits (0 - 9)
- **Task:** Create a model that can classify the digit (e.g. optical character recondition OCR)



dataset: <http://yann.lecun.com/exdb/mnist/>

Building a NN: Features

- Use the pixels of the image as features
- MNIST images are 28 x 28 so we will flatten to create a **784** dimensional feature vector.



A decorative graphic in the top-left corner consisting of several overlapping squares in shades of gray and purple.

Git Repo



<https://github.com/ericgossett/Intro-to-Neural-Networks-Tech-Talk>



Results

